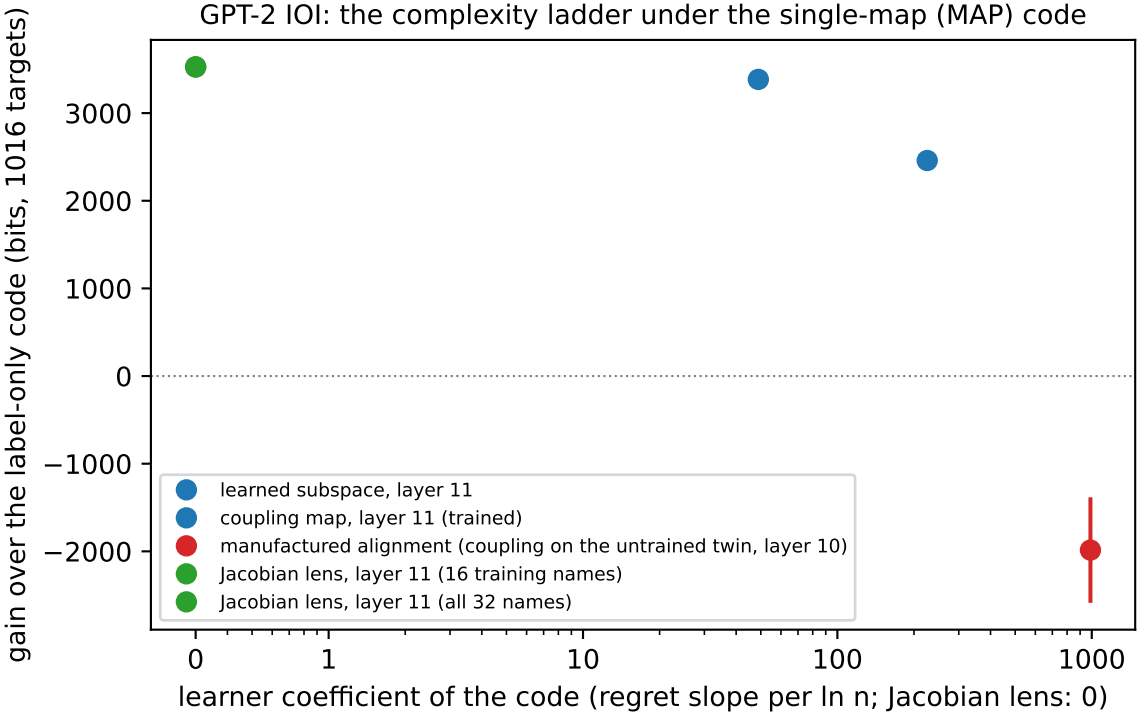
GPT-2 IOI: the complexity ladder under the single-map (MAP) code
gain over the label-only code (bits, 1016 targets)
3000
2000
1000
0
−1000
−2000
learned subspace, layer 11
coupling map, layer 11 (trained)
manufactured alignment (coupling on the untrained twin, layer 10)
Jacobian lens, layer 11 (16 training names)
Jacobian lens, layer 11 (all 32 names)
0
1
10
100
1000
learner coefficient of the code (regret slope per ln n; Jacobian lens: 0)